

Linear-time Suffix Array Construction by Induced Sorting, SAIS algorithm

Definition of a suffix. Given a text $T[0...n]$ over alphabet Σ of size σ , where the last character $\$$ of T is lexicographically smaller than any other character and occurs in T only once. A **suffix** T_i of T is a substring of $T[i...n]$.

Definition. A suffix T_i of T is called L-type if it is lexicographically greater than the suffix T_{i+1} , we denote this relation as $T_i > T_{i+1}$, and S-type if $T_i < T_{i+1}$ (T_i is lexicographically smaller than T_{i+1}). A suffix T_i is called leftmost S-type, LMS, if it is S-type, and suffix T_{i-1} is L-type.

Definition. A substring $T[i...j]$ of T is called LMS-Substring, or in short T^* substring, if T_i and T_j are LMS-suffixes, and none of the suffixes starting between i and j are LMS-suffix. We categorize a character $T[i]$ of T as L-type, S-type and LMS-type according to the category of T_i suffix.

Definition. A suffix array of T , SA , is an array $SA[0...n]$ such that $SA[i] = p$, where p is a position of suffix T_p and i is the lexicographic order of T_p among all suffixes of T . In other words, SA stores the positions of suffixes of T sorted in lexicographic (alphabetic) order.

SA-IS Algorithm Outline [1]

Input: Text $T[0...n]$

Output: Suffix Array of T , SA

0. Assign type of each suffix in T (L-type, S-type).
1. Induce-sort all LMS-substrings of T in $O(n)$ time.
2. Give each LMS-substring of T a name and construct a shortened string T_1 whose alphabet consists of the names of LMS-substrings. This step takes $O(n)$ time.
3. Call recursively SA-IS on T_1 to calculate the suffix array SA_1 for T_1 .
4. Induce SA from SA_1 (i.e. using SA_1 for T_1 , construct SA for T) in $O(n)$ time.

Time analysis of SA-IS:

The length of T_1 is at most $n/2$ since each character of T_1 corresponds to a T^* substring in T , and there are at most $n/2$ LMS-substrings in T (each LMS-substring starts with S-type character and the previous character in T is L-type character by definition). Thus the size of the original problem reduces in half with each recursive call. Time required to execute this algorithm is $T(n) = T(n/2) + O(n)$, which is $O(n)$.

Now we will describe how to implement each step of the algorithm.

We will use an example to illustrate all the steps of the algorithm. In our example,
 $T = \text{"m m i s s i s s i i p p i i \$"}'$

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$

Maintain an array $SA[0...n]$ (for suffix array), and initially $SA[i] = -1$, for all $0 \leq i \leq n$. Since array SA store positions of suffixes sorted in alphabetical order, suffixes starting with the same character span consecutive block of suffixes in SA (hereafter, we'll call a block of suffixes starting with the same character a **bucket**).

We'll need arrays A and C each of size equal to alphabet size σ . Array A at index c holds the total number of characters c in T. For example, in our example, $A[m] = 2$ and $A[i] = 6$. A can be constructed by scanning T once and counting the occurrence of each character c in T. Array C at index c holds the starting index of the c -bucket in SA. $C[c] = \sum_{i < c} Count(i)$, where $Count(i)$ is the number of all characters $i < c$ in T. In other words, $C[c]$ is the total number of characters in T alphabetically smaller than character c . In our example, $C[i] = 1$ (only one character \$ is smaller than i, so i-bucket will start at index 1 in SA) and $C[m] = 7$ (there is one character \$ and six characters i smaller than m , so m-bucket in SA starts at index 7).

For example, let character \$ be index 0, i be index 1, m be 2, p denote index 3 and s denote index 4 in A. Then after counting the number of occurrences of each character in T, A is

A =

Index of character c	$0_{\$}$	1_i	2_m	3_p	4_s
Occurrences of c in T	1	6	2	2	4

Then we scan A from index $j = 1, 2, 3, 4$ and calculate $Count(c)$, the total number of characters i in T that are lexicographically smaller than c , updating the values stored in C. $C[0] = 0$ (no characters are smaller than character \$ indexed by 0). $C[c] = C[c-1] + A[c-1]$.

C =

Index of character c	$0_{\$}$	1_i	2_m	3_p	4_s
$\sum_{i < c} Count(i)$	0	1	7	9	11

We need to maintain the pointers to heads (or ends) of c -buckets (a bucket starting with c -character) during the algorithm. So, we will keep another array B of size σ such that $B[c] = Head(c)$, where $Head(c)$ is the index of SA of the current head of the c -bucket. For Step 0, we need the information about current tails of c -buckets, i.e. pointers are placed at the **End** of each c -bucket.

So, initially, $B[c]$ keeps the index of the end of each c -bucket: $B[s] = n$, and $B[c] = B[c+1] - A[c+1]$.

B =

Index of character c	$0_{\$}$	1_i	2_m	3_p	4_s
End of c -bucket	0	6	8	10	14

Step 0. Assign type of each suffix in T (L-type, S-type). For this step, we'll need a bit-array (or bool array) t , such that $t[i] = 0$ if T_i is L-type suffix and $t[i] = 1$ if T_i is S-type suffix. Scan T from **Right-to-Left**. Initialize $\$$ as S-type, then if $T[i] > T[i+1]$ then $t[i] = 0$ (L-type), and if $T[i] < T[i+1]$ then $t[i] = 1$ (S-type). Finally if $T[i] = T[i+1]$, then check $t[i+1]$, if it is L-type, then $t[i]$ is also L-type; otherwise, $t[i]$ is S-type (see **Proof 4**).

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
t	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S

At the same time as assigning L-type to $T[i]$, check if $T[i+1]$ is S-type, then $T[i+1]$ is LMS, so store its position ($i+1$) into SA at the end of the $T[i+1]$ -character bucket:
 $SA[B[T[i+1]]] = i + 1$

For example, as we scan $T[13]$, we define it as L-type, check that $t[14]$ is S-type, and this means that $T[14]$ is LMS, so we need to put position 14 at the end of the $\$$ -bucket (where $\$$ is the character at position 14 in T). We find from $B[\$]$ that the end of this bucket is index 0 of SA, so put 14 into $SA[0]$ and update the current end of bucket by decrementing it by 1, i.e. $B[T[i+1]] = B[T[i+1]] - 1$. Now $B[\$] = 0 - 1 = -1$.

As we scan L-type $T[7]$, $t[8]$ is S-type, meaning $T[8]$ is LMS, and $T[8]$ is character i , and the current end of the i -bucket is $B[i] = 6$, so place 8 into $SA[6]$ and update the end of i -th bucket by decrementing 6 by 1, i.e. $B[i] = 6 - 1 = 5$.

After scanning T from right to left, we stored starting positions of all LMS-substrings at the end of corresponding c-buckets (where c is the starting character of an LMS-substring).

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14				2	5	8								

Step 1. Sort all LMS-substrings of T in $O(n)$ time. At the beginning of this step all LMS-substrings are placed in their corresponding buckets in SA (by Step 0).

a) Recalculate B array so that each $B[c]$ stores the **Head** of the c -bucket, i.e. $B[c] = C[c]$.

Index of character c	\$	i	m	p	s
$B[c]$	0	1	7	9	11

b) Induce-sort L-type suffixes using LMS-substrings, i.e. fill in SA[i] entries such that SA[i] = p and p is the starting position of an L-type suffix.

Scan SA from **Left-to-Right** and for each $p = SA[i]$ (such that p is not equal to -1), if T[p-1] is L-type (check if $t[p-1] = 0$), then store p-1 at the current head of the T[p-1]-bucket: $SA[B[T[p-1]]] = p-1$; and update $B[T[p-1]]$ by incrementing its value by 1, i.e. $B[T[p-1]] = B[T[p-1]] + 1$.

For example, starting at SA[0], $p = SA[0] = 14$, and $t[14-1] = 0$ (T[13] is L-type), so place $14-1 = 13$ into the current head of T[13] = i bucket. The head is located at $B[i] = 1$, so set $SA[1] = 13$. Update $B[i] = B[i] + 1 = 1 + 1 = 2$.

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12		2	5	8								

Next, $SA[1] = 13$, and $t[13-1]$ is L-type, so find the current head of T[12] = i bucket, $B[i] = 2$, and store 12 at SA[2]. Update $B[i] = 2 + 1 = 3$.

At the end of scanning of SA from left to right, SA is:

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12		2	5	8	1	0	11	10	4	7	3	6

c) Reset the values of B to point to the **End** of c-buckets. $B[s] = 14$, and $B[c] = B[c+1] - A[c+1]$.

d) Induce the order of S-type suffixes from ordered L-type suffixes, i.e. fill in entry SA[i] such that SA[i] = p and p is the starting position of S-type suffix.

Scan SA from **Right-to-Left**, and for each $SA[i] = p$ (valid, not -1), if T[p-1] is S-type suffix ($t[p-1] = 1$), then store p-1 at the end of the T[p-1]-bucket:

$SA[B[T[p-1]]] = p-1$;

and update the end of the bucket by moving it one to the left (decrement by 1): $B[T[p-1]] = B[T[p-1]] - 1$.

NOTE (for calculations of p-1 and p-2, previous two positions). While scanning SA, if $p = SA[i]$, and you need to calculate (p-1) or (p-2), there might occur out-of-range issue that you need to address appropriately. If $SA[i] = p = 0$, then p-1 is out-of-range index, use previous position to be equal to T.size - 1 (the last index) instead of $p - 1$. Previous position of position 0, is T.size - 1 (in our example it is 14), and instead of $p - 2$, use T.size - 2. Similarly, if position $p = 1$, then p-2 is out-of-range, so use T.size - 1 for this case (think of T as a circle starting at 0 and ending with T.size-1, so previous index of 0 on the circle is T.size-1).

For example, $SA[14] = 6$, and $T[6-1] = T[5]$ is S-type suffix, so place 5 at the end of T[5] = i bucket, which corresponds to index 6 of SA (note, SA[6] is over-written and now SA[6] = 5). Update $B[i] = 6 - 1 = 5$. Now the end of i-bucket is index 5 of SA.

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14

SA	14	13	12		2	5	5	1	0	11	10	4	7	3	6
----	----	----	----	--	---	---	---	---	---	----	----	---	---	---	---

Next, $SA[13] = 3$, and $T[3-1] = T[2]$ is S-type suffix, so place 2 at the end of $T[2]=i$ bucket, which corresponds to index 5 of SA (thus, $SA[5] = 2$). Update the end of the i -bucket by decrementing it by one: $B[i] = 5 - 1 = 4$.

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12		2	2	5	1	0	11	10	4	7	3	6

At the end of scanning from right to left, all LMS-substrings are correctly sorted relative to each other, but the order of S-type suffixes or L-type suffixes are not necessarily correct.

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
t	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12	8	9	2	5	1	0	11	10	4	7	3	6
L	1	0	0	1	0	1	1	0	0	0	0	0	0	0	0

At this step, we also maintain a bit-array (or a bool array) L such that $L[i] = 1$ if $SA[i] = p$ corresponds to LMS-substring starting at p , and $L[i] = 0$, otherwise. While scanning $SA[i] = p$ and checking whether $T[p-1]$ is of S-type, if so, then we also need to check whether $T[p-2]$ is of L-type, then $T[p-1]$ is LMS-substring, so set $L[B[T[p-1]]] = 1$.

At the end of Step 1, all LMS-substrings are correctly sorted relative to each other.

Step 2. Give each LMS-substring of T a name and construct a shortened string T_1 whose alphabet consists of the integer-names of LMS-substrings.

Idea is to scan SA from **Left-to-Right** and assign LMS-substrings names, integer-ranks of the LMS order in SA. For example if T_i is LMS-substring that occurs before another LMS-substring T_j in SA, then T_i 's name will be an integer less than T_j 's name if both substrings are different, or the same integer if the substrings are the same (same corresponding characters and same type characters, L-type character are smaller than S-type). Only two consecutive LMS-substrings in SA are needed to be compared (they are stored in alphabetic order, and only consecutive substrings may be the same). Write a function that compares two LMS-substrings (pass T by constant reference, pass array t , and two LMS substrings are passed each by an integer, starting position of the LMS in T); the function returns *true* if the two LMS substrings at position i and position j are same (same characters and same character-types) and returns *false* if the LMS substrings are different. Note that we can use information stored at **array L** to check if index i corresponds to $p = SA[i]$ such that $L[i] = 1$ (LMS starts at position p in T), but we don't know the length of an LMS. We can use a definition of LMS: the first and the last characters of LMS are LMS-characters and no other characters between the first and the last are LMS-characters; so as we scan t

array we can keep a track of whether the current character of either LMS substrings is LMS character, and if so, stop scanning (the end of one of LMS-substrings is reached).

The new shortened string T_1 is formed from the names (assigned integers) of LMS substrings preserving the order of occurrence of LMS-substrings in the original string T .

To clarify, the names are assigned 0, 1, 2, ..., to the LMS-substrings in the order of their occurrence in SA, but these names are put into a new string in the order of the corresponding LMS-substrings in the original string T .

We'll keep T_1 as an array of integers of size equal to the total number of LMS-substrings in T , call this number x (T_1 has alphabet $\{0, 1, \dots, k\}$, where $(k+1)$ is the size of the alphabet equal to the number of distinct LMS-substrings).

Now we'll explain how to implement this step. We need a temporarily pointer (variable) to one LMS-substring, that was processed immediately before LMS-substring that is being currently processed (call this pointer *previous*). Initially, *previous* is set to $SA[0]$:

$previous = SA[0]$

(a) The first LMS-substring corresponds to the special \$ character (sentinel), and it's name is 0 (no other LMS is identical to this one consisting of a single character). Set $T_1[x-1] = 0$ (the last character of T_1 is 0, the smallest in the alphabet).

(b) **Fill in N.** Maintain an array N (names of LMS substrings) of the same size as size of T (initialize $N[i] = -1$ for all $i = 0, 1, \dots, n$). $N[n] = 0$, (corresponds to the last character of the string \$).

Scan SA from **Left-to-Right**, and for each $SA[i]$, $i = 1 \dots n$, check if $L[i] = 1$. If so, then LMS-substring occurs at position $p = SA[i]$ in T , and this is currently processed LMS-substring. To assign a name to this LMS-substring, compare this current LMS to the LMS pointed by *previous* pointer. If these LMS substrings are identical, then the current LMS substring is assigned the same name as the previous LMS-substring; otherwise, increment the name of the previous LMS-substring by one and assign this integer to the current LMS-substring. Once we assigned a name m to an LMS-substring, we know that it occurs at position $p = SA[i]$ in T , so set $N[p] = N[SA[i]] = m$. Set *previous* to $SA[i]$.

(c) **Fill in T_1 .** Do this step only after array N is filled in (step b). To build T_1 , scan array N from **Left-to-Right** and fill in T_1 from **Left-to-Right**:

$T_1[j] = N[p]$ (where $N[p]$ is not -1)

Note, that we need to use two different variables j and p to access T_1 and N (because T_1 and N are of different lengths).

Example on the execution of Step 2.

c-bucket	\$	<i>i</i>						<i>m</i>		<i>p</i>		<i>s</i>			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12	8	9	2	5	1	0	11	10	4	7	3	6
L	1	0	0	1	0	1	1	0	0	0	0	0	0	0	0

While scanning SA from **Left-to-Right**, for $i = 1, \dots, n$ ($i = 0$ is processed already), we discover that $SA[3]$ corresponds to LMS-substring. Compare it to the previous LMS pointed by $previous = SA[0] = 14$. The current LMS-substring starting at position $SA[3] = 8$ in T is “iippii\$”, and the previous LMS-substring starting at position 14 in T is “\$”. The LMS-substrings are different, so the name for “iippii\$” is 1 (one greater than the previous name 0). Set $N[SA[3]] = N[8] = 1$ (the name of the LMS substring is 1, LMS substring starts at position 8 in T).

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
<i>t</i>	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S
<i>N</i>									1						0

The next LMS-substring is found at $SA[5]$, corresponding to position $SA[5] = 2$ in T, and the content of the current LMS-substring is now “issi”, which is not identical to the previous LMS-substring “iippii\$”, so the name of the “issi” is 2 (one greater than the name of the previous LMS-substring, which was 1). Set $N[SA[5]] = N[2] = 2$.

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
<i>t</i>	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S
<i>N</i>			2						1						0

The next LMS-substring is found at $SA[6]$ and it starts at position 5 in T corresponding to “issi”, it is identical to the previous LMS-substring (that started at position 2 of T), so the name of the current LMS-substring is the same as the name of the previous LMS-substring, namely 2. Set $N[SA[6]] = N[5] = 2$:

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
<i>t</i>	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S
<i>N</i>			2			2			1						0

Not filled entries in array N are initialized to -1. There are no more LMS-substrings. So now, we are ready to scan array N and fill in array T_1 :

Index of T_1	0	1	2	3
T_1	2	2	1	0

T_1 is now “2210”. Make a recursive call on T_1 .

Step 3. Call recursively SA-IS on T_1 to calculate the suffix array SA_1 for T_1 .

Termination step: check if all characters of T_1 are different (if the size of the alphabet is smaller than the length of T_1 , then some characters are repeated).

If all characters are different, we can calculate SA_1 by placing each character into its proper bucket in SA_1 and terminate the recursive call. For this, scan T_1 from **Left-to-Right** using index $i = 0, 1, 2, \dots$, and do:
 $SA_1[T_1[i]] = i$;

Step 4. After Recursive Call: Induce SA from SA_1 (i.e. using SA_1 for T_1 , construct SA for T) in $O(n)$ time.

The idea is to place sorted LMS-substrings in the same relative order from SA_1 to SA, and then induce-sort L-type suffixes from the order of LMS-substrings, and finally, induce-sort S-type suffixes from the order of L-type suffixes.

(a) Reset B so that $B[c]$ points to the **End** of c -bucket.

(b) Reset $SA[i] = -1$, for all $i = 0, 1, \dots, n$.

(c) Recall that i -th character in T_1 corresponds to i -th LMS-substring in the original string T .
 $SA_1[j] = i$ gives us the i -th suffix of T_1 that corresponds to i -th LMS-substring in T .

For this step, given i -th suffix of T_1 (hence, i -th LMS in T), we need to retrieve the starting position of i -th LMS in T . Since we do not need to use T_1 anymore, we may overwrite its i -th integer (corresponding to i -th LMS in T) with the starting position of i -th LMS in T . Scan t array (that holds L- and S-types) Left-to-Right using index j , and at the same time fill in T_1 using index i . If $t[j]$ is of S-type and $t[j-1]$ is of L-type, then $T[j]$ corresponds to LMS-substring, so set $T_1[i] = j$.

Now we can place positions of LMS-substrings into SA. Scan SA_1 from **Right-to-Left** for $j = x-1, \dots, 0$ (where x is the total number of LMS-substrings in T). We scan from right to left, because we want to put positions of LMS-substrings into SA at the end of the corresponding c -buckets (so all LMS-substrings are at the end of c -buckets in the preserved order of their occurrence in SA_1).

For each $SA_1[j] = i$, find $p = T_1[i]$, which is the starting position of the i -th LMS-substring in T . Then $T[p]$ is the character c , so we know in which c -bucket to place position p into SA.

$SA[B[T[p]]] = p$;

$B[T[p]] = B[T[p]] - 1$; decrement by 1 to point to the current not occupied end of the bucket.

Example of execution of Step 4.

Let SA_1 of T_1 be:

Index	0	1	2	3
T_1	2	2	1	0
SA_1	3	2	1	0

Overwrite T_1 with the starting positions of LMS-substrings in T using t -array:

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
t	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S

Index	0	1	2	3
T_1	2	5	8	14
SA_1	3	2	1	0

By scanning SA_1 from **Right-to-Left**, we find $SA_1[3] = 0$, which corresponds to position 0 in T_1 , and, hence, to the 0th LMS substring in T . Find position $p = T_1[0] = 2$. Find character $T[p] = T[2] = i$. Place position $p = 2$ into the end of i -bucket in SA , i.e. $SA[6] = 2$.

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA							2								

Next, $SA_1[2] = 1$, which corresponds to position 1 in T_1 , and to 1st LMS substring in T (count of LMS-substrings starts with 0). Find position $p = T_1[1] = 5$. Character $T[5] = i$, so place position 5 into the current end of the i -bucket.

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA						5	2								

Next, $SA_1[1] = 2$, which corresponds to position 2 in T_1 . Find $p = T_1[2] = 8$, and $T[8] = i$. Place 8 at the end of the i -bucket in SA .

c-bucket	\$	i						m		p		s			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA					8	5	2								

Finally, $SA_1[0] = 3$, corresponding to position 3 in T_1 . Find $p = T_1[3] = 14$. $T[14] = \$$, so place 14 into the end of $\$$ -bucket in SA :

c-bucket	\$	<i>i</i>						<i>m</i>		<i>p</i>		<i>s</i>			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14				8	5	2								

Now, all LMS-substrings's positions are placed into SA in correct relative order as they appear in SA₁.

Repeat induce-sort of L-type and S-type suffixes that we described in detail in Step 1 to obtain the final SA of T.

index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
T	m	m	i	s	s	i	s	s	i	i	p	p	i	i	\$
t	L	L	S	L	L	S	L	L	S	S	L	L	L	L	S

Reset B[c] to point to the **Head** of *c-bucket*.

B[c] = C[c].

Induce sort L-type suffixes:

c-bucket	\$	<i>i</i>						<i>m</i>		<i>p</i>		<i>s</i>			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12		8	5	2	1	0	11	10	7	4	6	3

Reset B[c] to point to the **End** of *c-bucket*.

B[s] = n, and B[c] = B[c+1] – A[c+1].

Induce sort S-type suffixes:

c-bucket	\$	<i>i</i>						<i>m</i>		<i>p</i>		<i>s</i>			
index	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
SA	14	13	12	8	9	5	2	1	0	11	10	7	4	6	3

Return SA.

Reference:

Ge Nong, Sen Zhang, and Wai Hong Chan, "Linear Suffix Array Construction by Almost Pure Induced-Sorting", *Data Compression Conference*, 2009, DCC '09, 16-18 March 2009, DOI: 10.1109/DCC.2009.42